

```
12s !pip install numpy pandas yfinance datetime scikit-learn

Requirement already satisfied: numpy in /usr/local/lib/python3.10/dist-packages (1.25.2)
Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (1.5.3)
Requirement already satisfied: yfinance in /usr/local/lib/python3.10/dist-packages (0.2.37)
Collecting datetime
  Downloading DateTime-5.4-py3-none-any.whl (52 kB)
    52.5/52.5 kB 1.7 MB/s eta 0:00:00
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.2.2)
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2.8.2)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas) (2023.4)
Requirement already satisfied: requests>=2.31 in /usr/local/lib/python3.10/dist-packages (from yfinance) (2.31.0)
Requirement already satisfied: multitasking>=0.0.7 in /usr/local/lib/python3.10/dist-packages (from yfinance) (0.0.11)
Requirement already satisfied: lxml>=4.9.1 in /usr/local/lib/python3.10/dist-packages (from yfinance) (4.9.4)
Requirement already satisfied: appdirs>=1.4.4 in /usr/local/lib/python3.10/dist-packages (from yfinance) (1.4.4)
Requirement already satisfied: frozendict>=2.3.4 in /usr/local/lib/python3.10/dist-packages (from yfinance) (2.4.0)
Requirement already satisfied: peewee>=3.16.2 in /usr/local/lib/python3.10/dist-packages (from yfinance) (3.17.1)
Requirement already satisfied: beautifulsoup4>=4.11.1 in /usr/local/lib/python3.10/dist-packages (from yfinance) (4.12.3)
Requirement already satisfied: html5lib>=1.1 in /usr/local/lib/python3.10/dist-packages (from yfinance) (1.1)
Collecting zope.interface (from datetime)
  Downloading zope.interface-6.2-cp310-cp310-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_17_x86_64.manylinux2014_x86_64.whl (247 kB)
    247.3/247.3 kB 7.4 MB/s eta 0:00:00
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.11.4)
Requirement already satisfied: joblib>=1.1.1 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.3.0)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4>=4.11.1->yfinance) (2.5)
Requirement already satisfied: six>=1.9 in /usr/local/lib/python3.10/dist-packages (from html5lib>=1.1->yfinance) (1.16.0)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from html5lib>=1.1->yfinance) (0.5.1)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests>=2.31->yfinance) (3.3.2)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests>=2.31->yfinance) (3.6)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests>=2.31->yfinance) (2.0.7)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests>=2.31->yfinance) (2024.2.2)
Requirement already satisfied: setuptools in /usr/local/lib/python3.10/dist-packages (from zope.interface->datetime) (67.7.2)
Installing collected packages: zope.interface, datetime
Successfully installed datetime-5.4 zope.interface-6.2
```

```
1s [2] # Task B2
import numpy as np
import pandas as pd
import yfinance as yf
import datetime as dt
from sklearn import preprocessing
from sklearn.model_selection import train_test_split
import os
```

We import the necessary modules and libraries for data processing and machine learning.

Import yfinance as yf: This line imports the yfinance module, which is a Python library that provides access to financial data from Yahoo Finance. It allows us to download the stock data for a given ticker and time period.

```
[3] # Task B2
TICKER = "AMZN"
START_DATE = "2014-01-01"
END_DATE = "2020-12-31"
LOOK_UP_DAYS = 30
TRAINING_RATIO = 0.9 # 0.7 == 70%
SCALE_DATA = True
SCALING_METHOD = "MinMax" # MinMax, Standard
```

We define all necessary constants and parameters for the data loading and splitting, such as the ticker symbol, the start and end dates, the look up days, the training ratio, the scale data flag, and the scaling method.

```
def DataScaler(stock_data, scaling_method=SCALING_METHOD):  
  
    DatasetScaler = None  
    ColumnScalers = {  
  
    }  
    if scaling_method == "MinMax":  
        DatasetScaler = preprocessing.MinMaxScaler()  
  
    elif scaling_method == "Standard":  
        DatasetScaler = preprocessing.StandardScaler()  
  
    # Learn the whole dataset  
    col_names = stock_data.columns  
    features = stock_data[col_names]  
    DatasetScaler.fit(features.values)  
    features = DatasetScaler.transform(features.values)  
    scaledDataFrame = pd.DataFrame(features, columns = col_names)  
    scaledDataFrame.index = stock_data.index  
  
    for column in col_names:  
        column_scaler = None  
        if scaling_method == "MinMax":  
            column_scaler = preprocessing.MinMaxScaler()  
        elif scaling_method == "Standard":  
            column_scaler = preprocessing.StandardScaler()  
        column_scaler.fit(stock_data[column].values.reshape(-1,1))  
        ColumnScalers[column] = column_scaler  
  
    return scaledDataFrame, DatasetScaler, ColumnScalers
```

Next, we define a function called DataScaler that takes a stock data DataFrame and a scaling method as inputs, and returns a scaled DataFrame, a dataset scaler, and a dictionary of column scalers. The scaling method can be either MinMax or Standard, and the function applies the scaling to each column of the stock data.

`DatasetScaler = preprocessing.MinMaxScaler()`: This line creates an instance of the `MinMaxScaler` class, which scales the data to be between 0 and 1. This is supposedly done to make the training process easier and faster, as the network can learn from smaller and normalized values.

`DatasetScaler = preprocessing.StandardScaler()`: This line creates an instance of the `StandardScaler` class, which standardizes the features in a dataset. Standardization means transforming the features to have zero mean and unit variance and can improve the performance of many machine learning algorithms. The `StandardScaler` class can compute the mean and standard deviation of the features on a training set, and then apply the same transformation on the testing set.

`column_scaler.fit(stock_data[column].values.reshape(-1,1))`: This line fits the `column_scaler` to the values of a specific column in the stock data. The `reshape` method changes the shape of the array to have one column and as many rows as needed. This is required because the scaler expects a two-dimensional input.

```
def DataLoader(ticker=TICKER, start_date=START_DATE, end_date=END_DATE, scale=SCALE_DATA, scaling_method=SCALING_METHOD):
    ...
    ticker: is the code of the target ticker
    start_date: a start date string with format YYYY/MM/DD
    end_date: an end date string with format YYYY/MM/DD
    scale: a boolean value, True by default
    scaling_method: MinMax(by default), Standard.
    ...

    # result
    result = {
        "dataset": None,
        "datasetScaler": None,
        "columnScalers": None
    }

    # processing the input parameters
    start_date = dt.datetime.strptime(start_date, "%Y-%m-%d")
    end_date = dt.datetime.strptime(end_date, "%Y-%m-%d")

    # creating necessary folder
    if not os.path.isdir("results"):
        os.mkdir("results")

    if not os.path.isdir("data"):
        os.mkdir("data")
```

```

# checking if the data is already downloaded
## Get a list of files in the directory
files = os.listdir("data")
## Check each file in the directory
data = None
for file_name in files:
    ## if we already downloaded the ticker data
    if file_name.startswith(ticker) and file_name.endswith(".csv"):
        ### Read the file
        file_path = os.path.join("data", f"{ticker}.csv")
        data = pd.read_csv(file_path, parse_dates=['Date'])
        break

## else, we gonna download the stock data
if data is None:
    stock_data = yf.download(ticker, start_date, end_date)
    file_path = os.path.join("data", f"{ticker}.csv")
    stock_data.to_csv(file_path)
    data = pd.read_csv(file_path, parse_dates=['Date'])

# if the given time is included in the file, we just take the nessecary dataframe
if data.head(1)["Date"].values[0] <= np.datetime64(start_date) and data.tail(1)["Date"].values[0] >= np.datetime64(end_date):
    data = data[(data['Date'] >= pd.to_datetime(start_date)) & (data['Date'] <= pd.to_datetime(end_date))]
    print("Local Stock Data is enough for requirements, do not need to download")
else:
    stock_data = yf.download(ticker, start_date, end_date)
    file_path = os.path.join("data", f"{ticker}.csv")
    stock_data.to_csv(file_path)
    data = pd.read_csv(file_path, parse_dates=['Date'])
    print("Local Stock Data is not enough for requirements, continuing downloading...")

# Setting Date as Index
data.set_index('Date', inplace=True)

```

```

# Scale Data
if scale:
    data, scaler, column_scalers = DataScaler(data, scaling_method)
    result["dataset"] = data
    result["datasetScaler"] = scaler
    result["columnScalers"] = column_scalers
    return result

result["dataset"] = data

return result

```

We define a function called DataLoader that takes a ticker, a start date, an end date, a scale flag, and a scaling method as inputs, and returns a result dictionary that contains the dataset, the dataset scaler, and the column scalers. The function checks if the data is already downloaded in the data folder, and if not, it downloads the data from yfinance. It also filters the data according to the given dates, and sets the date as the index. If the scale flag is True, it calls the DataScaler function defined above to scale the data.

stock_data = yf.download(ticker, start_date, end_date): This line downloads the stock data for the given ticker, start date, and end date using the yfinance module. It returns a pandas DataFrame with the date, open, high, low, , close, adjclose, and volume columns.

```
def datasetSplitter(dataset: pd.DataFrame, look_up_days=LOOK_UP_DAYS,
                    training_ratio=TRAINING_RATIO,
                    feature_columns=['Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume']):
    ...

    dataset: a Pandas Dataframe
    training_ratio: is equal to TRAINING_RATIO constant
    ...

    # result
    splitResult = {
        "X_training_set": {
            'Open': None,
            'High': None,
            'Low': None,
            'Close': None,
            'Adj Close': None,
            'Volume': None
        },
        "Y_training_set": {
            'Open': None,
            'High': None,
            'Low': None,
            'Close': None,
            'Adj Close': None,
            'Volume': None
        },
        "X_testing_set": {
            'Open': None,
            'High': None,
            'Low': None,
            'Close': None,
            'Adj Close': None,
            'Volume': None
        },
        "Y_testing_set": {
            'Open': None,
            'High': None,
            'Low': None,
            'Close': None,
            'Adj Close': None,
            'Volume': None
        }
    }
```

We then define a function called `datasetSplitter` that takes a dataset `DataFrame`, a look up days parameter, a training ratio parameter, and a list of feature columns as inputs, and returns a `splitResult` dictionary that contains the X and Y training and testing sets for each feature column. The function splits the data into X and Y arrays based on the look up days, and then splits them into training and testing sets based on the training ratio. It also converts the arrays to numpy arrays.

At the beginning of this function definition, we precisely specify the `splittedResult` dictionary to later correctly transfer the splitted data into their assigned label.

```
for column in feature_columns:
    dataset_in_column = dataset[column].values.reshape(-1, 1)    # <class 'numpy.ndarray'>
    x_data = []
    y_data = []

    for x in range(look_up_days, len(dataset_in_column)):
        x_data.append(dataset_in_column[x - look_up_days:x, 0])
        y_data.append(dataset_in_column[x, 0])

    splitResult["X_training_set"][column], splitResult["X_testing_set"][column], splitResult["Y_training_set"][column], splitResult["Y_testing_set"][column] = train_test_split(x_data, y_data, test_size=0.2, random_state=42)

    ## Converting to numpy.array

    for column in feature_columns:
        splitResult["X_training_set"][column] = np.array(splitResult["X_training_set"][column])
        splitResult["Y_training_set"][column] = np.array(splitResult["Y_training_set"][column])
        splitResult["X_testing_set"][column] = np.array(splitResult["X_testing_set"][column])
        splitResult["Y_testing_set"][column] = np.array(splitResult["Y_testing_set"][column])

return splitResult
```

The loop here extracts the values as a numpy array and reshapes it to have one column and as many rows as needed. This is done to make the array compatible with the scaler and the neural network model.

`x_data.append(dataset_in_column[x - look_up_days:x, 0])`: This line appends a slice of the `dataset_in_column` input array to the `x_data` list. The slice contains the values from the `x - look_up_days` row to the `x-1` row, and from the 0th column.

`y_data.append(dataset_in_column[x, 0])`: This line appends the value of the `dataset_in_column` output array to the `y_data` list. The value is at the `x`th row and the 0th column. This means that for each time `t`, the `y_data` will contain the value of the column at time `t`. For example, if `t` is 31, then `y_data` will contain the value at row 31, which is the value at day 31.

```

def Data_Processing_1():
    StockData = DataLoader()

    scaledStockData, datasetScaler, columnScalers = DataScaler(stock_data=StockData["dataset"])

    dataset = datasetSplitter(dataset=scaledStockData)

    print("Loaded Done!\nThe result is a tuple as below:\n")
    print("(dataset, scaledStockData, datasetScaler, columnScalers), where:\n")
    print("dataset is a dictionary as below:")

```

```

print("(dataset, scaledStockData, datasetScaler, columnScalers), where:\n")
print("dataset is a dictionary as below:")
print('''{
    "X_training_set": {
        'Open': <class 'numpy.ndarray'>,
        'High': <class 'numpy.ndarray'>,
        'Low': <class 'numpy.ndarray'>,
        'Close': <class 'numpy.ndarray'>,
        'Adj Close': <class 'numpy.ndarray'>,
        'Volume': <class 'numpy.ndarray'>
    },
    "Y_training_set": {
        'Open': ...,
        ...
    },
    "X_testing_set": {
        'Open': ...,
        'High': ...,
        ...
    },
    "Y_testing_set": {
        'Open': ...,
        'High': ...,
        'Low': ...,
        ...
    }
}\n''')
print("scaledStockData is a Pandas Dataframe of the Stock Ticker (scaled)\n")
print("datasetScaler is the Scaler of the dataset\n")
print("columnScalers is a dictionary: each key is a DataFrame Feature('Open', 'High', etc.) and the correspond value is a scaler of that feature")

return dataset, scaledStockData, datasetScaler, columnScalers

```

We defines a function called `Data_Processing_1` that calls the `DataLoader` function to get the stock data, the `DataScaler` function to scale the data, and the `datasetSplitter` function to split the data. It also prints some fundamental information about the result tuple, which contains the dataset, the scaledStockData, the datasetScaler, and the columnScalers.

```
dataset, scaledStockData, datasetScaler, ColumnScalers = Data_Processing_1()

[*****100%*****] 1 of 1 completed
[*****100%*****] 1 of 1 completedLocal Stock Data is not enough for requirements, continuing downloading...
Loaded Done!
The result is a tuple as below:

(dataset, scaledStockData, datasetScaler, columnScalers), where:

dataset is a dictionary as below:
{
    "X_training_set": {
        'Open': <class 'numpy.ndarray'>,
        'High': <class 'numpy.ndarray'>,
        'Low': <class 'numpy.ndarray'>,
        'Close': <class 'numpy.ndarray'>,
        'Adj Close': <class 'numpy.ndarray'>,
        'Volume': <class 'numpy.ndarray'>
    },
    "Y_training_set": {
        'Open': ...,
        ...
    },
    "X_testing_set": {
        'Open': ...,
        'High': ...,
        ...
    },
    "Y_testing_set": {
        'Open': ...,
        'High': ...,
        'Low': ...,
        ...
    }
}

scaledStockData is a Pandas Dataframe of the Stock Ticker (scaled)

datasetScaler is the Scaler of the dataset

columnScalers is a dictionary: each key is a DataFrame Feature('Open', 'High', etc.) and the correspond value is a scaler of that feature
```

And here is the results from the Data_Processing_1() function, with all required output described in the Task specification.

We also consider to improve our current data preparation model by implementing the sklearn.model_selection.TimeSeriesSplit method to split our data into training and testing sets in a time-series manner since this will preserve the temporal order and structure of our data, and avoid any data leakage or contamination. Besides, we can also use the sklearn.metrics module to evaluate our model performance on the testing set, such as the mean absolute error, the mean squared error, the root mean squared error, the R-squared score, etc.

```
from sklearn.model_selection import TimeSeriesSplit

def datasetSplitter(dataset: pd.DataFrame, look_up_days=LOOK_UP_DAYS, training_ratio=TRAINING_RATIO, feature_columns=['Open', 'High', 'Low', 'Close', 'Adj Close', 'Volume']):
    ...
    dataset: a Pandas Dataframe
    training_ratio: is equal to TRAINING_RATIO constant
    ...
```



```
#result:
splitResult = {
  "X_training_set": {
    "Open": None,
    "High": None,
    "Low": None,
    "Close": None,
    "Adj Close": None,
    "Volume": None
  },
  "Y_training_set": {
    "Open": None,
    "High": None,
    "Low": None,
    "Close": None,
    "Adj Close": None,
    "Volume": None
  },
  "X_testing_set": {
    "Open": None,
    "High": None,
    "Low": None,
    "Close": None,
    "Adj Close": None,
    "Volume": None
  },
  "Y_testing_set": {
    "Open": None,
    "High": None,
    "Low": None,
    "Close": None,
    "Adj Close": None,
    "Volume": None
  }
}
```



✓
1s




```
scaledStockData.plot()
```



```
<Axes: xlabel='Date'>
```



 print(dataset)

```
{ 'X_training_set': { 'Open': array([[0.03506406, 0.03490774, 0.03415987, ..., 0.02373566, 0.02207749,
    0.0194017 ],
    [0.03490774, 0.03415987, 0.03391161, ..., 0.02207749, 0.0194017 ,
    0.0229694 ],
    [0.03415987, 0.03391161, 0.03496291, ..., 0.0194017 , 0.0229694 ,
    0.021725 ],
    ...,
    [0.45652546, 0.48626246, 0.48226874, ..., 0.63188867, 0.63995891,
    0.64535949],
    [0.48626246, 0.48226874, 0.44062402, ..., 0.63995891, 0.64535949,
    0.65353089],
    [0.48226874, 0.44062402, 0.45074479, ..., 0.64535949, 0.65353089,
    0.63893823]]), 'High': array([[0.03339843, 0.03442546, 0.03267491, ..., 0.02243832, 0.02064179,
    0.02047317],
    [0.03442546, 0.03267491, 0.03312558, ..., 0.02064179, 0.02047317,
    0.02112924],
    [0.03267491, 0.03312558, 0.03451437, ..., 0.02047317, 0.02112924,
    0.02002251],
    ...,
    [0.48204532, 0.49170251, 0.48466659, ..., 0.66544854, 0.64674737,
    0.66053717],
    [0.49170251, 0.48466659, 0.45207138, ..., 0.64674737, 0.66053717,
    0.65542652],
    [0.48466659, 0.45207138, 0.45860451, ..., 0.66053717, 0.65542652,
    0.64490789]]), 'Low': array([[0.03435237, 0.0350393 , 0.03260384, ..., 0.02255916, 0.01945864,
    0.01951485],
    [0.0350393 , 0.03260384, 0.03443668, ..., 0.01945864, 0.01951485,
    0.02165367],
    [0.03260384, 0.03443668, 0.03498309, ..., 0.01951485, 0.02165367,
    0.02043595],
    ...,
    [0.46126536, 0.47902544, 0.47382042, ..., 0.64039915, 0.6344729 ,
    0.65633886],
    [0.47902544, 0.47382042, 0.43432239, ..., 0.6344729 , 0.65633886,
    0.62311994],
    [0.47382042, 0.43432239, 0.43607715, ..., 0.65633886, 0.62311994,
    0.64539497]]), 'Close': array([[0.03421791, 0.03374635, 0.03288026, ..., 0.02306673, 0.01920173,
    0.02165203],
    [0.03374635, 0.03288026, 0.03423641, ..., 0.01920173, 0.02165203,
```



print(dataset)



```
0.14498059]])), 'Y_training_set': {'Open': array([0.0229694 , 0.021725 , 0.02091583, ..., 0.65353089, 0.63893823,
0.64843374]), 'High': array([0.02112924, 0.02002251, 0.01965767, ..., 0.65542652, 0.64490789,
0.65417264]), 'Low': array([0.02165367, 0.02043595, 0.01938995, ..., 0.62311994, 0.64539497,
0.65509927]), 'Close': array([0.02169826, 0.02055787, 0.01862536, ..., 0.62911698, 0.64001851,
0.65110191]), 'Adj Close': array([0.02169826, 0.02055787, 0.01862536, ..., 0.62911698, 0.64001851,
0.65110191]), 'Volume': array([0.11486063, 0.17918328, 0.14306109, ..., 0.28707105, 0.14498059,
0.18216916])}), 'X_testing_set': {'Open': array([[0.44062402, 0.45074479, 0.41595965, ..., 0.65353089, 0.63893823,
0.64843374],
[0.45074479, 0.41595965, 0.45701891, ..., 0.63893823, 0.64843374,
0.65365043],
[0.41595965, 0.45701891, 0.44921227, ..., 0.64843374, 0.65365043,
0.66168087],
...,
[0.88136757, 0.8697358 , 0.86090847, ..., 0.89451353, 0.89517559,
0.89177344],
[0.8697358 , 0.86090847, 0.88859803, ..., 0.89517559, 0.89177344,
0.89180403],
[0.86090847, 0.88859803, 0.87341379, ..., 0.89177344, 0.89180403,
0.92734009]]), 'High': array([[0.45207138, 0.45860451, 0.4503699 , ..., 0.65542652, 0.64490789,
0.65417264],
[0.45860451, 0.4503699 , 0.48051553, ..., 0.64490789, 0.65417264,
0.65301072],
[0.4503699 , 0.48051553, 0.47557352, ..., 0.65417264, 0.65301072,
0.66050653],
...,
[0.88461387, 0.87414118, 0.8744416 , ..., 0.89875317, 0.89511408,
0.89262165],
[0.87414118, 0.8744416 , 0.88871272, ..., 0.89511408, 0.89262165,
0.92389238],
[0.8744416 , 0.88871272, 0.87361388, ..., 0.89262165, 0.92389238,
0.93819419]]), 'Low': array([[0.43432239, 0.43607715, 0.41903211, ..., 0.62311994, 0.64539497,
0.65509927],
[0.43607715, 0.41903211, 0.43876864, ..., 0.64539497, 0.65509927,
0.65507431],
[0.41903211, 0.43876864, 0.45617902, ..., 0.65509927, 0.65507431,
0.64914181],
...,
[0.87490517, 0.87469907, 0.87073369, ..., 0.90426482, 0.90554186,
0.90080522],
```

✓
0s



scaledStockData.head(10)



	Open	High	Low	Close	Adj Close	Volume
Date						
2014-01-02	0.035064	0.033398	0.034352	0.034218	0.034218	0.054690
2014-01-03	0.034908	0.034425	0.035039	0.033746	0.033746	0.057842
2014-01-06	0.034160	0.032675	0.032604	0.032880	0.032880	0.099644
2014-01-07	0.033912	0.033126	0.034437	0.034236	0.034236	0.045036
2014-01-08	0.034963	0.034514	0.034983	0.035435	0.035435	0.062468
2014-01-09	0.036569	0.035707	0.035732	0.035155	0.035155	0.053176
2014-01-10	0.036207	0.034747	0.034284	0.034122	0.034122	0.078268
2014-01-13	0.034813	0.033527	0.032613	0.032063	0.032063	0.085468
2014-01-14	0.033020	0.033175	0.033500	0.034085	0.034085	0.063496
2014-01-15	0.035107	0.033383	0.033887	0.033571	0.033571	0.078216

